

Evaluation

September 12, 2025

1 Reading in the dataframe

```
[6]: import pandas as pd
import sklearn
import numpy as np
import matplotlib.pyplot as plt
import krippendorff
from sklearn.metrics import classification_report

sample = "Set-1" #Or Set-2
if sample == "Set-1":
    file = "/home/pc/Uni/MasterThesis/Scripts/
↳Corpus_abortion_sample3_chunked_annotated.csv"
    file2 = "/home/pc/Uni/MasterThesis/Scripts/
↳Corpus_econ_sample_chunked_annotated.csv"

elif sample == "Set-2":
    file = "/home/pc/Uni/MasterThesis/Scripts/
↳Corpus_abort_sample2_chunked_annotated.csv"
    file2 = "/home/pc/Uni/MasterThesis/Scripts/
↳Corpus_econ_sample2_chunked_annotated.csv"

Corpus_abortion_sample_chunked_annotated_cleaned = pd.read_csv(file)
Corpus_abortion_sample_chunked_annotated_cleaned = ↳
    Corpus_abortion_sample_chunked_annotated_cleaned.fillna(99)
Corpus_abortion_sample_chunked_annotated_cleaned['expert'] = ↳
    Corpus_abortion_sample_chunked_annotated_cleaned['expert'].astype(float).
    astype('int')

Corpus_econ_sample_chunked_annotated_cleaned = pd.read_csv(file2)
Corpus_econ_sample_chunked_annotated_cleaned = ↳
    Corpus_econ_sample_chunked_annotated_cleaned.fillna(99)
Corpus_econ_sample_chunked_annotated_cleaned['expert'] = ↳
    Corpus_econ_sample_chunked_annotated_cleaned['expert'].astype(float).
    astype('int')

del file, file2
```

1.1 Standardizing dataframes

```
[2]: Column_ordering = ["chunk", "speaker", "file_name", "date", "expert", "student_assistant", "phi:2.7b", "gemma3:4b", "qwen2.5:7b", "mistral:7b", "llama3.2:latest", "stablelm2:12b"] #Standard ordering since the first annotation was less structured

Corpus_econ_sample_chunked_annotated_cleaned = Corpus_econ_sample_chunked_annotated_cleaned.reindex(columns=Column_ordering)
Corpus_abortion_sample_chunked_annotated_cleaned = Corpus_abortion_sample_chunked_annotated_cleaned.reindex(columns=Column_ordering)

if Corpus_abortion_sample_chunked_annotated_cleaned['student_assistant'].hasnans: #Exclude student assistant as the column is only featured in the first set of samples
    Corpus_abortion_sample_chunked_annotated_cleaned.drop(columns=['student_assistant'], inplace=True)
if Corpus_econ_sample_chunked_annotated_cleaned['student_assistant'].hasnans:
    Corpus_econ_sample_chunked_annotated_cleaned.drop(columns=['student_assistant'], inplace=True)

if 'Column_ordering' in locals():
    del Column_ordering
```

2 Evaluating metrics

2.1 Speed metrics

```
[3]: performance_econ_df = pd.read_csv("Model_metrics_econ_sample.csv")
performance_econ_df = performance_econ_df.sort_values(by='Speed', ascending=False)
performance_econ_df = performance_econ_df.reset_index(drop=True)

performance_econ_df['Speed'] = performance_econ_df['Speed'] / performance_econ_df['Speed'].max() #Standardizing the speed metric by the fastest model
print(performance_econ_df[['Model', 'Speed']])
```

	Model	Speed
0	phi:2.7b	1.000000
1	llama3.2:latest	0.988633
2	gemma3:4b	0.951229
3	qwen2.5:7b	0.891019
4	mistral:7b	0.839343
5	stablelm2:12b	0.752199

```
[4]: performance_abortion_df = pd.read_csv("Model_metrics_abortion_sample.csv")
performance_abortion_df = performance_abortion_df.sort_values(by='Speed', ↴
    ascending=False)
performance_abortion_df = performance_abortion_df.reset_index(drop=True)

performance_abortion_df['Speed'] = performance_abortion_df['Speed'] / ↴
    performance_abortion_df['Speed'].max()
print(performance_abortion_df[['Model', 'Speed']])
```

	Model	Speed
0	phi:2.7b	1.000000
1	llama3.2:latest	0.976086
2	gemma3:4b	0.942400
3	mistral:7b	0.877817
4	qwen2.5:7b	0.865682
5	stablelm2:12b	0.784476

```
[5]: def model_quality_metrics(sample):
    report = pd.DataFrame(columns=['Model', 'Scores'])

    for model_name in sample.columns[5:]:
        print(f"Model: {model_name}")
        sample[model_name] = sample[model_name].astype(int)
        print(classification_report(sample['expert'], sample[model_name], ↴
            labels=[1, 0], target_names=["Topic mentioned", "No mentions"], ↴
            zero_division=np.nan))

        report.loc[len(report)] = {'Model': model_name, 'Scores': ↴
            classification_report(sample['expert'], sample[model_name], ↴
            output_dict=True, zero_division=np.nan)}

    return report

Report_abortion_classification = ↴
    model_quality_metrics(Corpus_abortion_sample_chunked_annotated_cleaned)
Report_economic_classification = ↴
    model_quality_metrics(Corpus_econ_sample_chunked_annotated_cleaned)
```

	precision	recall	f1-score	support
Topic mentioned	0.69	0.64	0.67	64
No mentions	0.75	0.79	0.77	86
accuracy			0.73	150
macro avg	0.72	0.72	0.72	150
weighted avg	0.72	0.73	0.72	150

Model: gemma3:4b

	precision	recall	f1-score	support
Topic mentioned	0.81	0.98	0.89	64
No mentions	0.99	0.83	0.90	86
accuracy			0.89	150
macro avg	0.90	0.90	0.89	150
weighted avg	0.91	0.89	0.89	150

Model: qwen2.5:7b

	precision	recall	f1-score	support
Topic mentioned	0.93	0.89	0.91	64
No mentions	0.92	0.95	0.94	86
accuracy			0.93	150
macro avg	0.93	0.92	0.92	150
weighted avg	0.93	0.93	0.93	150

Model: mistral:7b

	precision	recall	f1-score	support
Topic mentioned	nan	0.00	0.00	64
No mentions	0.57	1.00	0.73	86
accuracy			0.57	150
macro avg	0.57	0.50	0.36	150
weighted avg	0.57	0.57	0.42	150

Model: llama3.2:latest

	precision	recall	f1-score	support
Topic mentioned	0.89	0.80	0.84	64
No mentions	0.86	0.93	0.89	86
accuracy			0.87	150
macro avg	0.88	0.86	0.87	150
weighted avg	0.87	0.87	0.87	150

Model: stablelm2:12b

	precision	recall	f1-score	support
Topic mentioned	0.95	0.86	0.90	64
No mentions	0.90	0.97	0.93	86
accuracy			0.92	150
macro avg	0.93	0.91	0.92	150

weighted avg	0.92	0.92	0.92	150
Model: student assistant				
	precision	recall	f1-score	support
Topic mentioned	0.67	0.94	0.78	70
No mentions	0.92	0.59	0.72	80
accuracy			0.75	150
macro avg	0.79	0.77	0.75	150
weighted avg	0.80	0.75	0.75	150
Model: phi:2.7b				
	precision	recall	f1-score	support
Topic mentioned	0.51	0.63	0.56	70
No mentions	0.59	0.46	0.52	80
accuracy			0.54	150
macro avg	0.55	0.55	0.54	150
weighted avg	0.55	0.54	0.54	150
Model: gemma3:4b				
	precision	recall	f1-score	support
Topic mentioned	0.56	0.94	0.70	70
No mentions	0.88	0.35	0.50	80
accuracy			0.63	150
macro avg	0.72	0.65	0.60	150
weighted avg	0.73	0.63	0.59	150
Model: qwen2.5:7b				
	precision	recall	f1-score	support
Topic mentioned	0.95	0.59	0.73	70
No mentions	0.73	0.97	0.83	80
accuracy			0.79	150
macro avg	0.84	0.78	0.78	150
weighted avg	0.83	0.79	0.78	150
Model: mistral:7b				
	precision	recall	f1-score	support
Topic mentioned	0.84	0.90	0.87	70
No mentions	0.91	0.85	0.88	80

accuracy		0.87	150	
macro avg	0.87	0.88	0.87	150
weighted avg	0.88	0.87	0.87	150

Model: llama3.2:latest

	precision	recall	f1-score	support
Topic mentioned	0.94	0.73	0.82	70
No mentions	0.80	0.96	0.88	80

accuracy		0.85	150	
macro avg	0.87	0.85	0.85	150
weighted avg	0.87	0.85	0.85	150

Model: stablelm2:12b

	precision	recall	f1-score	support
Topic mentioned	0.91	0.73	0.81	70
No mentions	0.80	0.94	0.86	80

accuracy		0.84	150	
macro avg	0.85	0.83	0.84	150
weighted avg	0.85	0.84	0.84	150

```
[ ]: def create_classification_heatmaps(report_abortion, report_economic, sample):
    import seaborn as sns
    import numpy as np

    def extract_metrics(report_df, corpus_name):
        metrics_data = []
        for _, row in report_df.iterrows():
            model_name = row['Model']
            scores = row['Scores']

            metrics_data.append({
                'Model': model_name,
                'Accuracy': scores['accuracy'],
                'F1_Class_1': scores['1']['f1-score'],
                'F1_Class_0': scores['0']['f1-score'],
                'Precision_Class_1': scores['1']['precision'],
                'Recall_Class_1': scores['1']['recall'],
                'Precision_Class_0': scores['0']['precision'],
                'Recall_Class_0': scores['0']['recall']
            })
        return pd.DataFrame(metrics_data)
```

```

if sample == "Set-1":
    sample_desc = 'Abortion annotation task evaluation metrics heatmap'
    ↪\n(Sample 1: 1974)'
    sample_desc2 = 'Economic annotation task evaluation metrics heatmap'
    ↪\n(Sample 1: random draw)'
elif sample == "Set-2":
    sample_desc = 'Abortion annotation task evaluation metrics heatmap'
    ↪\n(Sample 2: 2022)'
    sample_desc2 = 'Economic annotation task evaluation metrics heatmap'
    ↪\n(Sample 2: random draw post 2020)'

abortion_metrics = extract_metrics(report_abortion, 'Abortion annotation'
task')
economic_metrics = extract_metrics(report_economic, 'Economic annotation'
task')

fig, axes = plt.subplots(1, 2, figsize=(20, 8))

# Abortion heatmap
abortion_heatmap_data = abortion_metrics.set_index('Model')[['Accuracy', 'F1_Class_1', 'F1_Class_0', 'Precision_Class_1', 'Precision_Class_0', 'Recall_Class_1', 'Recall_Class_0']]
abortion_hm = sns.heatmap(
    abortion_heatmap_data.T, annot=True, cmap='RdYlGn', fmt='.3f',
    cbar_kws={'label': 'Score'}, ax=axes[0],
    annot_kws={"size": 10, "family": "DejaVu Serif"}, vmin=0.4, vmax=1)
cbar0 = abortion_hm.collections[0].colorbar
cbar0.ax.yaxis.label.set_fontproperties('DejaVu Serif')
cbar0.ax.yaxis.label.set_size(10)
cbar0.set_ticks([0.4, 0.5, 0.75, 1])
for label in cbar0.ax.get_yticklabels():
    label.set_fontname('DejaVu Serif')
    label.set_fontsize(10)
axes[0].set_title(sample_desc, fontsize=14, fontweight='bold', family='DejaVu Serif')
axes[0].set_xlabel('Models', fontsize=10, family='DejaVu Serif')
axes[0].set_xticklabels(axes[0].get_xticklabels(), rotation=45, horizontalalignment='right', family='DejaVu Serif', fontsize=10)
axes[0].set_yticklabels(axes[0].get_yticklabels(), rotation=0, family='DejaVu Serif', fontsize=10)

# Economic heatmap

```

```


economic_heatmap_data = economic_metrics.set_index('Model')[['Accuracy', □
↳ 'F1_Class_1', 'F1_Class_0', 'Precision_Class_1', 'Precision_Class_0', □
↳ 'Recall_Class_1', 'Recall_Class_0']]
economic_hm = sns.heatmap(
    economic_heatmap_data.T, annot=True, cmap='RdYlGn', fmt='.3f',
    cbar_kws={'label': 'Score'}, ax=axes[1],
    annot_kws={"size": 10, "family": "DejaVu Serif"}, □
    vmin=0.4, vmax=1
)
cbar1 = economic_hm.collections[0].colorbar
cbar1.ax.yaxis.label.set_fontproperties('DejaVu Serif')
cbar1.ax.yaxis.label.set_size(10)
cbar1.set_ticks([0.4, 0.5, 0.75, 1])
for label in cbar1.ax.get_yticklabels():
    label.set_fontname('DejaVu Serif')
    label.set_fontsize(10)
axes[1].set_title(sample_desc2, fontsize=14, fontweight='bold', □
↳ family='DejaVu Serif')
axes[1].set_xlabel('Models', fontsize=10, family='DejaVu Serif')
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=45, □
↳ horizontalalignment='right', family='DejaVu Serif', fontsize=10)
axes[1].set_yticklabels(axes[1].get_yticklabels(), rotation=0, □
↳ family='DejaVu Serif', fontsize=10)

plt.tight_layout()
plt.savefig('classification_metrics_heatmaps_sample1.png', dpi=300, □
↳ bbox_inches='tight')
plt.show()

return abortion_metrics, economic_metrics

```

abortion_data, economic_data =
↳ create_classification_heatmaps(Report_abortion_classification, □
↳ Report_economic_classification, sample)



3 Plotting differences

```
[8]: f1_scored_abort_df = pd.DataFrame(columns=['Model', 'F1_0', 'F1_1'])

for model_name in Corpus_abortion_sample_chunked_annotation_cleaned.columns[5:12]:
    Corpus_abortion_sample_chunked_annotation_cleaned[model_name] = Corpus_abortion_sample_chunked_annotation_cleaned[model_name].astype(int)
    report = classification_report(
        Corpus_abortion_sample_chunked_annotation_cleaned['expert'],
        Corpus_abortion_sample_chunked_annotation_cleaned[model_name],
        output_dict=True,
        zero_division=np.nan
    )

    f1_scored_abort_df.loc[len(f1_scored_abort_df)] = {'Model': model_name, 'F1_0': report['0']['f1-score'], 'F1_1': report['1']['f1-score']}

f1_scored_abort_df = f1_scored_abort_df.sort_values(by='F1_1', ascending=False).reset_index().drop(columns=['index'])

plt.figure(figsize=(14, 7))

bar_width = 0.35
x = np.arange(len(f1_scored_abort_df['Model']))

bar1 = plt.bar(x - bar_width/2, f1_scored_abort_df['F1_1'], width=bar_width, color='black')
```

```


bar2 = plt.bar(x + bar_width/2, f1_scored_abort_df['F1_0'], width=bar_width, color='grey')

plt.ylabel('F1 Score', fontsize=10, family='DejaVu Serif')
plt.bar_label(bar1, fontsize=10, fmt='%.3f')
plt.bar_label(bar2, fontsize=10, fmt='%.3f')
plt.title('F1 Scores by Model and Class for annotation of abortion topics', fontsize=12, fontweight='bold', family='DejaVu Serif')
plt.xlabel('Models', fontsize=10, family='DejaVu Serif')
plt.xticks(x, f1_scored_abort_df['Model'], rotation=45, ha='right', fontname='DejaVu Serif', fontsize=10)
plt.legend([bar1, bar2], ['Topic mentioned', 'No mentions'], prop={'family':'DejaVu Serif', 'size':10})
plt.yticks(fontsize=10)
plt.ylim(0, 1)

plt.tight_layout()
plt.style.use('fivethirtyeight')
plt.show()

def bar1, bar2, x, bar_width, report, model_name #Saving RAM saves lives.

```



```
[9]: plt.figure(figsize=(14, 6))
plt.bar(performance_abortion_df['Model'], height=performance_abortion_df['Speed'], color = 'darkgrey')
plt.xlabel('Models', fontsize=10, family='DejaVu Serif')
plt.ylabel('Speed standardized', fontsize=10, family='DejaVu Serif')
```

```

plt.title('Speed standardized by model for annotation of abortion topics',  

    fontsize=12, fontweight='bold', family='DejaVu Serif')  

plt.xticks(rotation=45, ha='right', fontname='DejaVu Serif', fontsize = 10)  

plt.yticks(np.arange(0, 1.1, 0.1), ['0%', '10%', '20%', '30%', '40%', '50%',  


    '60%', '70%', '80%', '90%', '100%'], fontsize=10)  

plt.style.use('fivethirtyeight')  

plt.savefig('abortion_annotation_speed.png', dpi=300, bbox_inches='tight')  

plt.show()

```



```

[10]: f1_scored_econ_df = pd.DataFrame(columns=['Model', 'F1_0', 'F1_1'])

for model_name in Corpus_econ_sample_chunked_annotated_cleaned.columns[5:12]:  

    report = classification_report(  

        Corpus_econ_sample_chunked_annotated_cleaned['expert'],  

        Corpus_econ_sample_chunked_annotated_cleaned[model_name],  

        output_dict=True  

    )  

    f1_scored_econ_df.loc[len(f1_scored_econ_df)] = {'Model': model_name,  

        'F1_0': report['0']['f1-score'], 'F1_1': report['1']['f1-score']}  

f1_scored_econ_df = f1_scored_econ_df.sort_values(by='F1_1', ascending=False).  

    reset_index().drop(columns=['index'])  

plt.figure(figsize=(14, 7))  

bar_width = 0.35  

x = np.arange(len(f1_scored_econ_df['Model']))

```

```


bar1 = plt.bar(x - bar_width/2, f1_scored_econ_df['F1_1'], width=bar_width, color='black')
bar2 = plt.bar(x + bar_width/2, f1_scored_econ_df['F1_0'], width=bar_width, color='grey')

plt.ylabel('F1 Score', fontsize=10, family='DejaVu Serif')
plt.bar_label(bar1, fontsize=10, fmt='%.3f')
plt.bar_label(bar2, fontsize=10, fmt='%.3f')
plt.title('F1 Scores by Model and Class for annotation of economic topics', fontsize=12, fontweight='bold', family='DejaVu Serif')
plt.xlabel('Models', fontsize=10, family='DejaVu Serif')
plt.xticks(x, f1_scored_econ_df['Model'], rotation=45, ha='right', fontname='DejaVu Serif', fontsize=10)
plt.legend([bar1, bar2], ['Topic mentioned', 'No mentions'], prop={'family': 'DejaVu Serif', 'size': 10})
plt.yticks(fontsize=10)
plt.ylim(0, 1)

plt.tight_layout()
plt.style.use('fivethirtyeight')
plt.show()

def bar1, bar2, x, bar_width, report, model_name

```



```


[11]: plt.figure(figsize=(14, 7))
plt.bar(performance_econ_df['Model'], height=performance_econ_df['Speed'], color = 'darkgrey')
plt.xlabel('Models', fontsize=10, family='DejaVu Serif')

```

```

plt.ylabel('Speed standardized', fontsize=10, family='DejaVu Serif')
plt.yticks(np.arange(0, 1.1, 0.1), ['0%', '10%', '20%', '30%', '40%', '50%', '60%', '70%', '80%', '90%', '100%'], fontsize=10)
plt.title('Speed standardized by model for annotation of economic topics', fontweight='bold', family='DejaVu Serif')
plt.xticks(rotation=45, ha='right', fontname='DejaVu Serif', fontsize=10)
plt.style.use('fivethirtyeight')
plt.savefig('economic_annotation_speed.png', dpi=300, bbox_inches='tight')
plt.show()

```



3.1 Excluding outlier models

```

[12]: Corpus_abortion_selected =
    Corpus_abortion_sample_chunked_annotated_cleaned[['qwen2.5:7b', 'gemma3:4b',
    "llama3.2:latest", "stablelm2:12b", "expert"]] #removing Phi and Mistral due
    to especially low performance
Corpus_econ_selected = Corpus_econ_sample_chunked_annotated_cleaned[['gemma3:
    4b', "llama3.2:latest", "stablelm2:12b", "mistral:7b", "expert"]] #removing
    Phi and Mistral due to especially low performance

```

3.2 Calculating intermodel-agreement

```
[13]: def krippendorff_alpha(corpus1, corpus2):
```

```

annotation_cols_abort = corpus1.columns[~corpus1.columns.
↪isin(['mean_agreement', 'agreement', 'expert', 'student_assistant', 'chunk', ↵
↪'speaker', 'date', 'file_name', 'processing_type'])]
annotation_cols_econ = corpus2.columns[~corpus2.columns.
↪isin(['mean_agreement', 'agreement', 'expert', 'student_assistant', 'chunk', ↵
↪'speaker', 'date', 'file_name', 'processing_type'])]

data_for_alpha_abort = corpus1[annotation_cols_abort].T.values
data_for_alpha_econ = corpus2[annotation_cols_econ].T.values

alpha_abort = krippendorff.alpha(reliability_data=data_for_alpha_abort, ↵
↪level_of_measurement='nominal')
print(f"Krippendorff's Alpha (Abortion Corpus): {alpha_abort:.3f}")

alpha_econ = krippendorff.alpha(reliability_data=data_for_alpha_econ, ↵
↪level_of_measurement='nominal')
print(f"Krippendorff's Alpha (Economic Corpus): {alpha_econ:.3f}")

def agreement_score(corpus1, corpus2, sigma=0.166):
    cols_to_average1 = corpus1.columns[~corpus1.columns.isin(['mean_agreement', ↵
↪'agreement', 'expert', 'student_assistant', 'chunk', 'speaker', 'date', ↵
↪'file_name', 'processing_type'])]
    cols_to_average2 = corpus2.columns[~corpus2.columns.isin(['mean_agreement', ↵
↪'agreement', 'expert', 'student_assistant', 'chunk', 'speaker', 'date', ↵
↪'file_name', 'processing_type'])]

    corpus1['mean_agreement'] = corpus1[cols_to_average1].mean(axis=1)
    corpus2['mean_agreement'] = corpus2[cols_to_average2].mean(axis=1)

    def inverted_gaussian(x, mu=0.5, sigma=sigma):
        return 1 - np.exp(-((x - mu) ** 2) / (2 * sigma ** 2))

    corpus1['agreement'] = corpus1['mean_agreement'].apply(inverted_gaussian)
    corpus2['agreement'] = corpus2['mean_agreement'].apply(inverted_gaussian)

    return corpus1, corpus2

if __name__ == "__main__":
    print("Krippendorff for selected models:\n")
    krippendorff_alpha(Corpus_abortion_selected, Corpus_econ_selected)
    print("\nKrippendorff for all models:\n")
    krippendorff_alpha(Corpus_abortion_sample_chunked_annotated_cleaned, ↵
↪Corpus_econ_sample_chunked_annotated_cleaned)

```

```

    agreement_score(Corpus_abortion_sample_chunked_annotated_cleaned, ↵
Corpus_econ_sample_chunked_annotated_cleaned)
    agreement_score(Corpus_abortion_selected, Corpus_econ_selected)

    print(f'\nAgreement metrics for the abortion sample')
    ↵\n{Corpus_abortion_sample_chunked_annotated_cleaned['agreement'].describe()}\n')
    print(f'Agreement metrics for the econ sample')
    ↵\n{Corpus_econ_sample_chunked_annotated_cleaned['agreement'].describe()}\n')

    print("Agreement score after removing Phi and Mistral \n")
    print(f'Agreement metrics for the abortion sample w.o. Phi and Mistral\n')
    ↵{Corpus_abortion_selected['agreement'].describe()}\n')
    print(f'Agreement metrics for the econ sample w.o. Phi')
    ↵{Corpus_econ_selected['agreement'].describe()}\n')

```

Krippendorff for selected models:

Krippendorff's Alpha (Abortion Corpus): 0.750
Krippendorff's Alpha (Economic Corpus): 0.399

Krippendorff for all models:

Krippendorff's Alpha (Abortion Corpus): 0.400
Krippendorff's Alpha (Economic Corpus): 0.311

Agreement metrics for the abortion sample

count	150.000000
mean	0.758623
std	0.301154
min	0.000000
25%	0.395905
50%	0.866826
75%	0.989286
max	0.989286
Name:	agreement, dtype: float64

Agreement metrics for the econ sample

count	150.000000
mean	0.624645
std	0.365436
min	0.088420
25%	0.088420
50%	0.901174
75%	0.901174
max	0.989286
Name:	agreement, dtype: float64

```
Agreement score after removing Phi and Mistral
```

```
Agreement metrics for the abortion sample w.o. Phi and Mistral
```

```
count    150.000000
```

```
mean     0.871123
```

```
std      0.269217
```

```
min     0.000000
```

```
25%     0.989286
```

```
50%     0.989286
```

```
75%     0.989286
```

```
max     0.989286
```

```
Name: agreement, dtype: float64
```

```
Agreement metrics for the econ sample w.o. Phi count    150.000000
```

```
mean     0.752189
```

```
std      0.276020
```

```
min     0.000000
```

```
25%     0.678274
```

```
50%     0.678274
```

```
75%     0.989286
```

```
max     0.989286
```

```
Name: agreement, dtype: float64
```

```
/tmp/ipykernel_74902/2519639514.py:19: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
corpus1['mean_agreement'] = corpus1[cols_to_average1].mean(axis=1)
```

```
/tmp/ipykernel_74902/2519639514.py:20: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
corpus2['mean_agreement'] = corpus2[cols_to_average2].mean(axis=1)
```

```
/tmp/ipykernel_74902/2519639514.py:25: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
corpus1['agreement'] = corpus1['mean_agreement'].apply(inverted_gaussian)
```

```
/tmp/ipykernel_74902/2519639514.py:26: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
corpus2['agreement'] = corpus2['mean_agreement'].apply(inverted_gaussian)
```

3.3 Evaluating problematic chunks

```
[14]: print(f'Mean length for problematic chunks (abortion):')
    ↪{Corpus_abortion_sample_chunked_annotated_cleaned[Corpus_abortion_sample_chunked_annotated_
    ↪<= 0.5] ['chunk'].apply(len).mean().round()} \n')
print(f'Mean length for problematic chunks (economic):'
    ↪{Corpus_econ_sample_chunked_annotated_cleaned[Corpus_econ_sample_chunked_annotated_cleaned[
    ↪<= 0.5] ['chunk'].apply(len).mean().round()} \n')

print(f'Mean length non-problematic chunks (abortion):'
    ↪{Corpus_abortion_sample_chunked_annotated_cleaned[Corpus_abortion_sample_chunked_annotated_
    ↪> 0.5] ['chunk'].apply(len).mean().round()} \n')
print(f'Mean length non-problematic chunks (economic):'
    ↪{Corpus_econ_sample_chunked_annotated_cleaned[Corpus_econ_sample_chunked_annotated_cleaned[
    ↪> 0.5] ['chunk'].apply(len).mean().round()} \n')
```

Mean length for problematic chunks (abortion): 6425.0

Mean length for problematic chunks (economic): 7117.0

Mean length non-problematic chunks (abortion): 2945.0

Mean length non-problematic chunks (economic): 4456.0